

RECALL DIFFERENT SORTING TECHNIQUES SO FAR

- ① Heap sort : $O(n \log n)$
 - ② Insertion sort
 - ③ Selection sort
 - ④ In-order traversal on binary search tree / AVL tree
 - Construction = $O(n^2)$
 - Traversal = $O(n)$
 - Construction = $O(n \log n)$
 - Traversal = $O(n)$
- ④ also gives you first or last K sorted elements

Priority Queue Sorting (§ 7.1.4)

- ◆ We can use a priority queue to sort a set of comparable elements
 1. Insert the elements one by one with a series of **insert** operations
 2. Remove the elements in sorted order with a series of **removeMin** operations
- ◆ The running time of this sorting method depends on the priority queue implementation

Algorithm **PQ-Sort(S, C)**

Input sequence S , comparator C for the elements of S

Output sequence S sorted in increasing order according to C

$P \leftarrow$ priority queue with comparator C

while $\neg S.isEmpty()$

$e \leftarrow S.removeFirst()$

$P.insert(e, 0)$

while $\neg P.isEmpty()$

$e \leftarrow P.removeMin().key()$

$S.insertLast(e)$

Selection-Sort

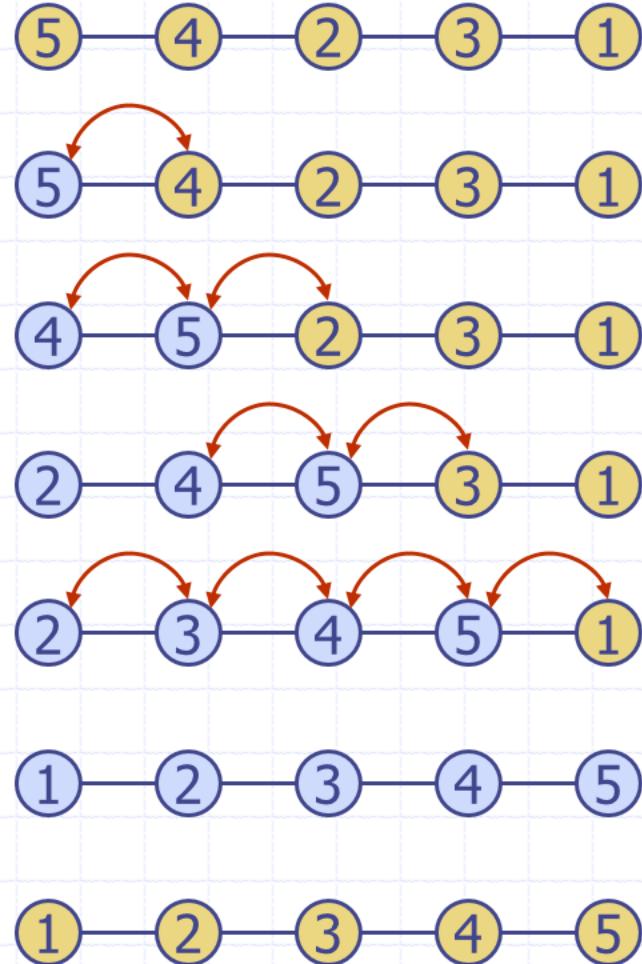
- ◆ Selection-sort is the variation of PQ-sort where the priority queue is implemented with an unsorted sequence
- ◆ Running time of Selection-sort:
 1. Inserting the elements into the priority queue with n insert operations takes $O(n)$ time
 2. Removing the elements in sorted order from the priority queue with n `removeMin` operations takes time proportional to
$$1 + 2 + \dots + n$$
- ◆ Selection-sort runs in $O(n^2)$ time

Insertion-Sort

- ◆ Insertion-sort is the variation of PQ-sort where the priority queue is implemented with a sorted sequence
- ◆ Running time of Insertion-sort:
 1. Inserting the elements into the priority queue with n **insert** operations takes time proportional to
$$1 + 2 + \dots + n$$
 2. Removing the elements in sorted order from the priority queue with a series of n **removeMin** operations takes $O(n)$ time
- ◆ Insertion-sort runs in $O(n^2)$ time

In-place Insertion-sort (w/o PQ)

- ◆ Instead of using an external data structure, we can implement selection-sort and insertion-sort in-place
- ◆ A portion of the input sequence itself serves as the priority queue
- ◆ For in-place insertion-sort
 - We keep sorted the initial portion of the sequence
 - We can use **swaps** instead of modifying the sequence



Heap-Sort (§2.4.4)

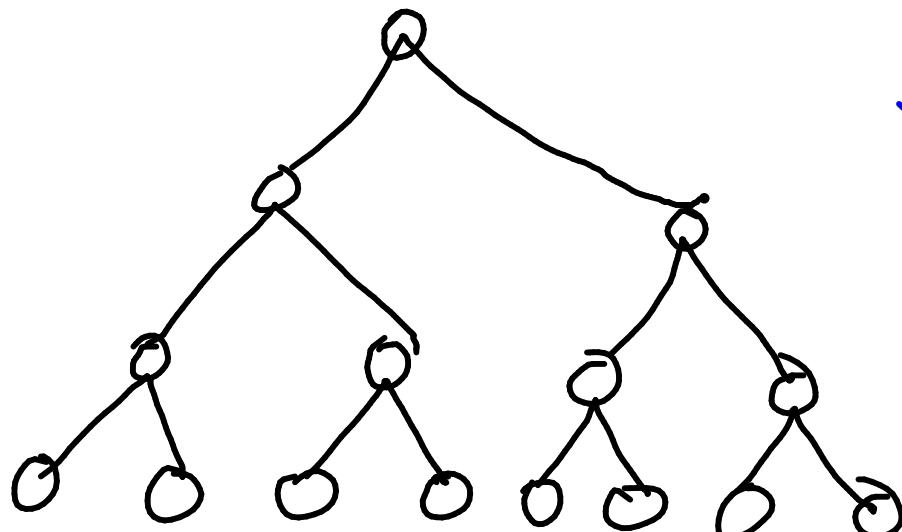


- ◆ Consider a priority queue with n items implemented by means of a heap
 - the space used is $O(n)$
 - methods `insert` and `removeMin` take $O(\log n)$ time
 - methods `size`, `isEmpty`, and `min` take time $O(1)$ time
- invoke removeMin n times*

- ◆ Using a heap-based priority queue, we can sort a sequence of n elements in $O(n \log n)$ time
- ◆ The resulting algorithm is called heap-sort
- ◆ Heap-sort is much faster than quadratic sorting algorithms, such as insertion-sort and selection-sort

H/w question: Understand why heap construction can be done in $O(n)$ time

Bottom-up heap construction



height # trees

$$\{ 3 : 2^0 \}$$

$$\{ 2 : 2^1 \}$$

$$\{ 1 : 2^2 \}$$

$$\{ 0 : 2^3 \}$$

$$\sum_{i=0}^h i \times 2^{h-i} = 2^h \left[\sum_{i=0}^h \frac{i}{2^i} \right]$$

$$O(2^n) = O(n)$$

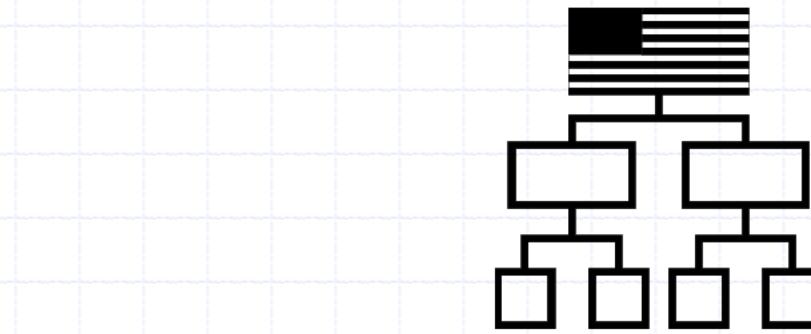
2^i nodes at level i

$$n = 2^{h+1} - 1$$

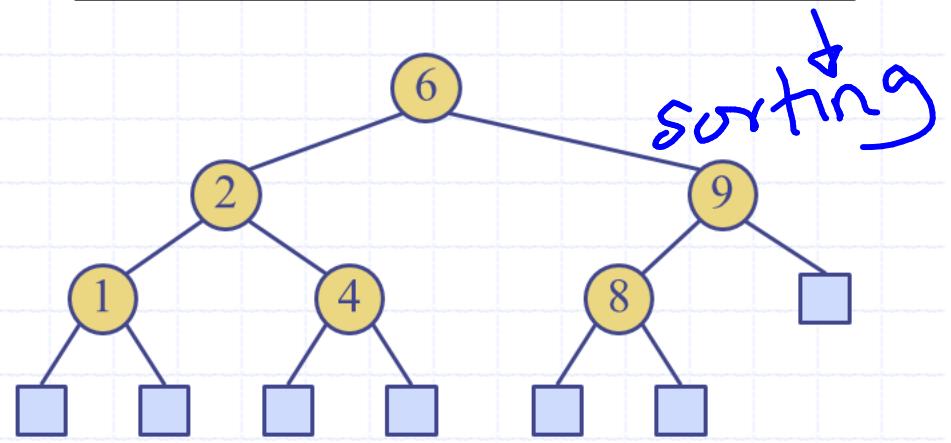
$$\sum_{i=0}^{\infty} \frac{1}{x^i} = \frac{x}{x-1} \quad \text{take deriv.}$$

Binary Search Trees (§ 9.1)

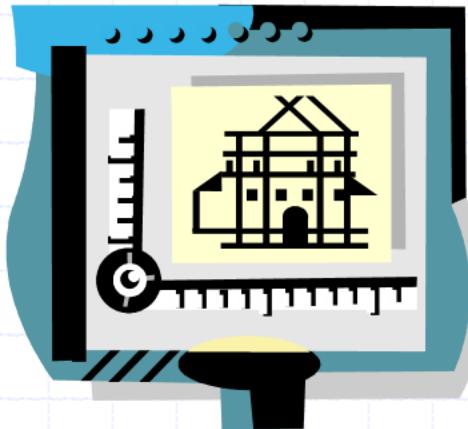
- ◆ A binary search tree is a binary tree storing keys (or key-value entries) at its internal nodes and satisfying the following property:
 - Let u , v , and w be three nodes such that u is in the left subtree of v and w is in the right subtree of v . We have $\text{key}(u) \leq \text{key}(v) \leq \text{key}(w)$
- ◆ External nodes do not store items



- ◆ An inorder traversal of a binary search trees visits the keys in increasing order



Running Times for AVL Trees



- ◆ a single restructure is $O(1)$
 - using a linked-structure binary tree
- ◆ find is $O(\log n)$
 - height of tree is $O(\log n)$, no restructures needed
- ◆ insert is $O(\log n)$
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$
- ◆ remove is $O(\log n)$
 - initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$

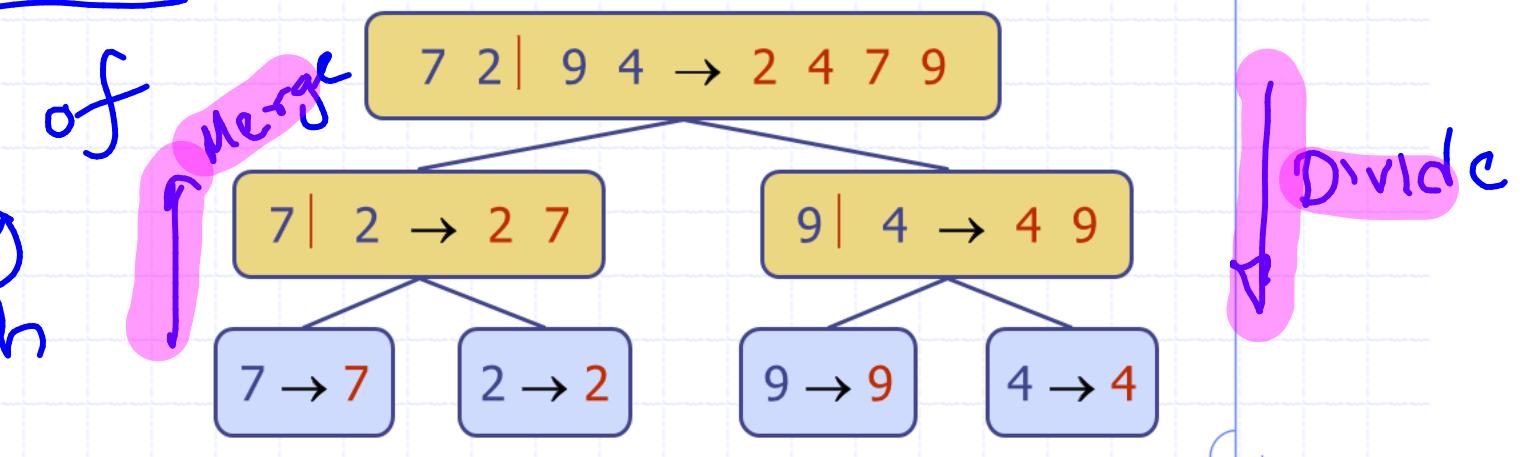
∴ Constructing an AVL tree of n elements

$$= n \times \log n^{11}$$

$$= O(n \log n)$$

Merge Sort

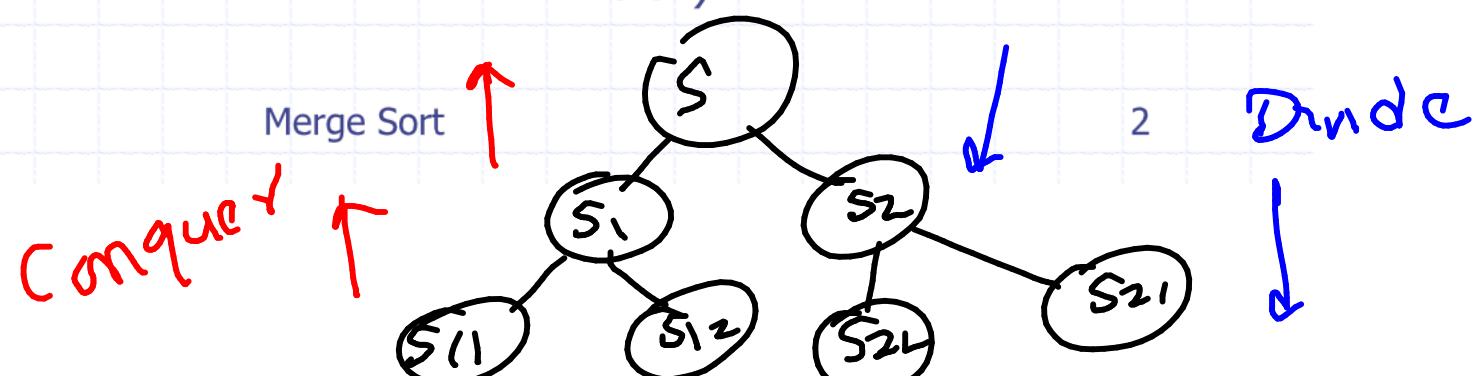
Generalise idea of
sorting using
binary search
trees



Divide-and-Conquer (§ 10.1.1)

- ◆ Divide-and conquer is a general algorithm design paradigm:
 - **Divide**: divide the input data S in two disjoint subsets S_1 and S_2
 - **Recur**: solve the subproblems associated with S_1 and S_2
 - **Conquer**: combine the solutions for S_1 and S_2 into a solution for S
 - ◆ The base case for the recursion are subproblems of size 0 or 1
- Sk < Ps* { } { }

- ◆ Merge-sort is a sorting algorithm based on the divide-and-conquer paradigm
- ◆ Like heap-sort
 - It uses a comparator
 - It has $O(n \log n)$ running time
- ◆ Unlike heap-sort
 - It does not use an auxiliary priority queue
 - It accesses data in a sequential manner (suitable to sort data on a disk)



Merge-Sort (§ 10.1)

- ◆ Merge-sort on an input sequence S with n elements consists of three steps:
 - **Divide**: partition S into two sequences S_1 and S_2 of about $n/2$ elements each
 - **Recur**: recursively sort S_1 and S_2
 - **Conquer**: merge S_1 and S_2 into a unique sorted sequence
- answering
load balancing
already sorted*

Algorithm *mergeSort(S, C)*

Input sequence S with n elements, comparator C

Output sequence S sorted according to C

if $S.size() > 1$

$(S_1, S_2) \leftarrow partition(S, n/2)$

mergeSort(S_1, C)

mergeSort(S_2, C)

$S \leftarrow merge(S_1, S_2)$

$S_i: [e_{i1} e_{i2} \dots e_{ik}]$ → ascending

$S_L: [e_{l1} e_{l2} \dots e_{lk}]$ → ascending

if $e_{i3} < e_{lk}$
 $\Rightarrow e_{i1} \neq e_{i2} < e_{lk}$
if $e_{ij} < e_{lk}$

$i < j, e_{ij} < e_{lk}$

Merging Two Sorted Sequences

- ◆ The conquer step of merge-sort consists of merging two sorted sequences A and B into a sorted sequence S containing the union of the elements of A and B
- ◆ Merging two sorted sequences, each with $n/2$ elements and implemented by means of a doubly linked list, takes $O(n)$ time

Algorithm $\text{merge}(A, B)$

Input sequences A and B with $n/2$ elements each

Output sorted sequence of $A \cup B$

$S \leftarrow$ empty sequence

while $\neg A.\text{isEmpty}() \wedge \neg B.\text{isEmpty}()$

if $A.\text{first}().\text{element}() < B.\text{first}().\text{element}()$
 $S.\text{insertLast}(A.\text{remove}(A.\text{first}()))$

else

$S.\text{insertLast}(B.\text{remove}(B.\text{first}()))$

while $\neg A.\text{isEmpty}()$

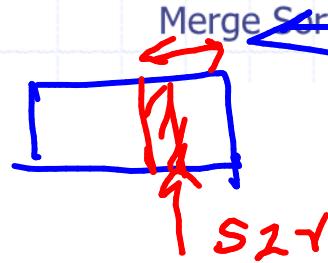
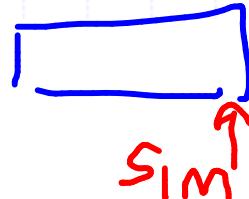
$S.\text{insertLast}(A.\text{remove}(A.\text{first}()))$

while $\neg B.\text{isEmpty}()$

$S.\text{insertLast}(B.\text{remove}(B.\text{first}()))$

return S

Not necessary
that A & B
have same
length

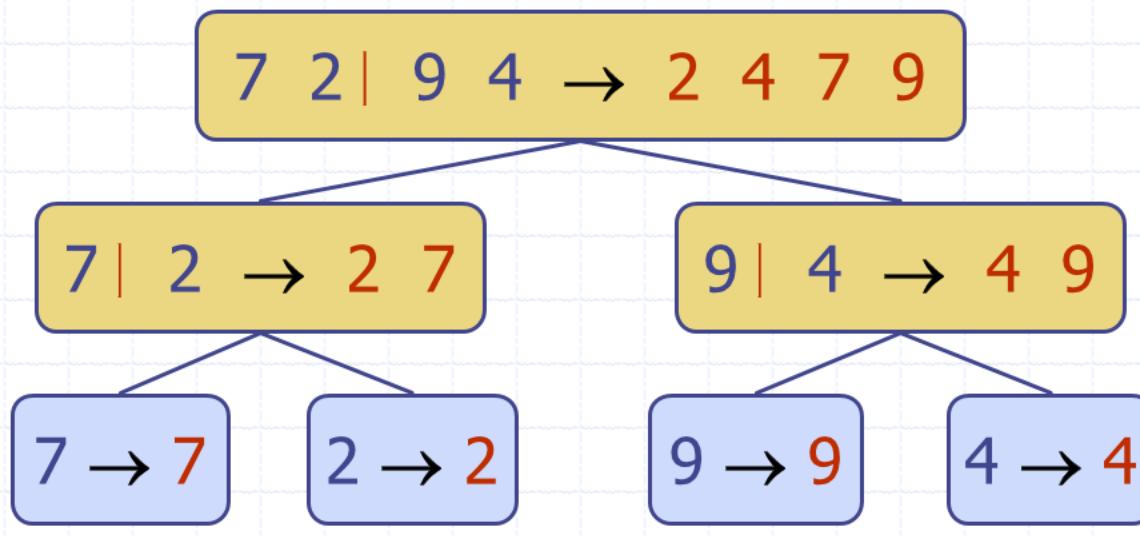


Merge Sort

Handling the excess
elements after all
comparisons over

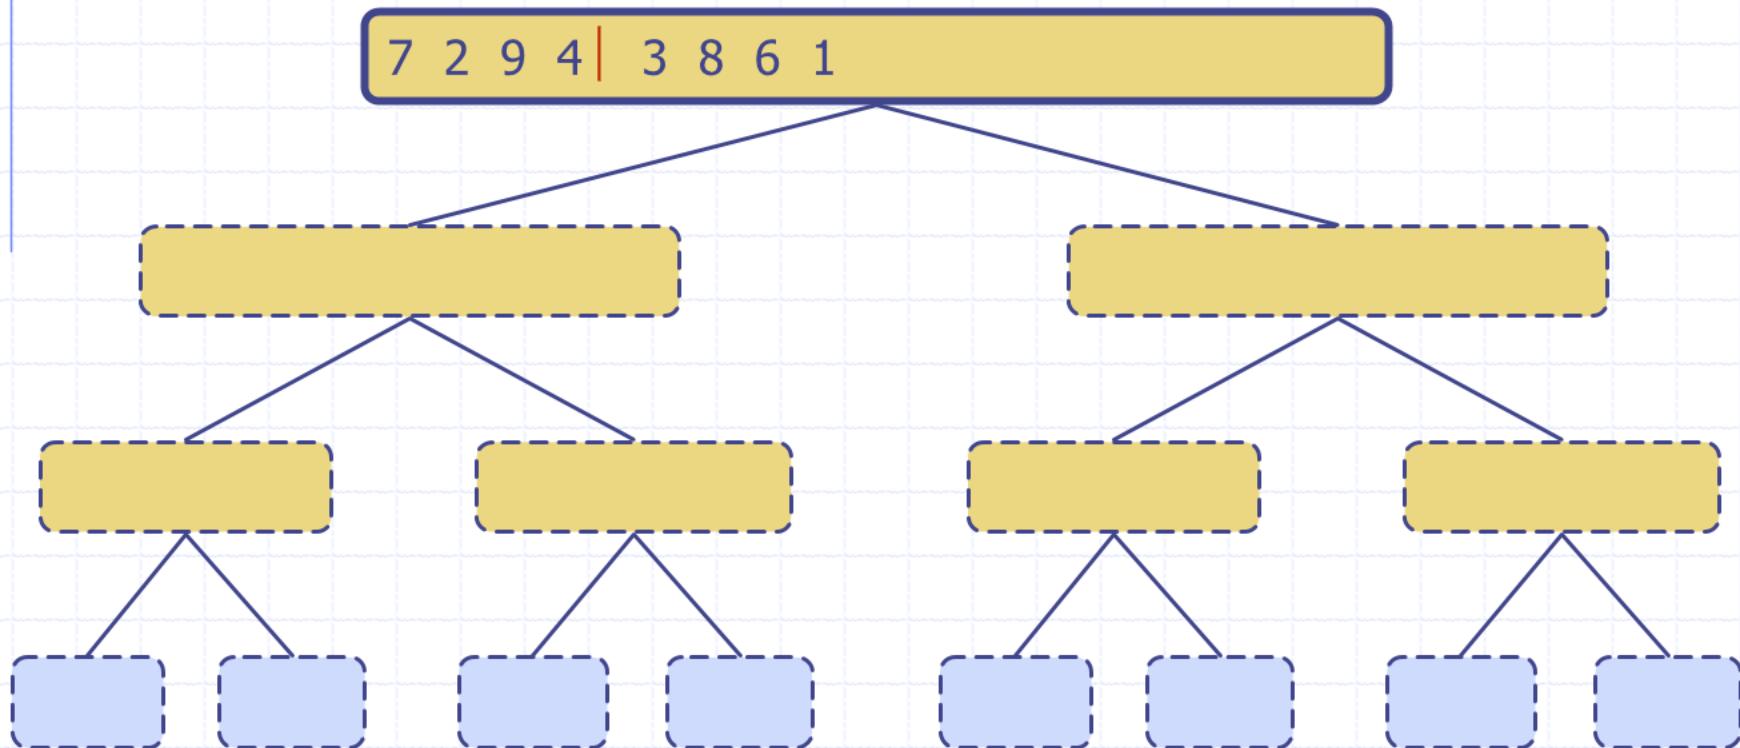
Merge-Sort Tree

- ◆ An execution of merge-sort is depicted by a binary tree
 - each node represents a recursive call of merge-sort and stores
 - ◆ unsorted sequence before the execution and its partition
 - ◆ sorted sequence at the end of the execution
 - the root is the initial call
 - the leaves are calls on subsequences of size 0 or 1



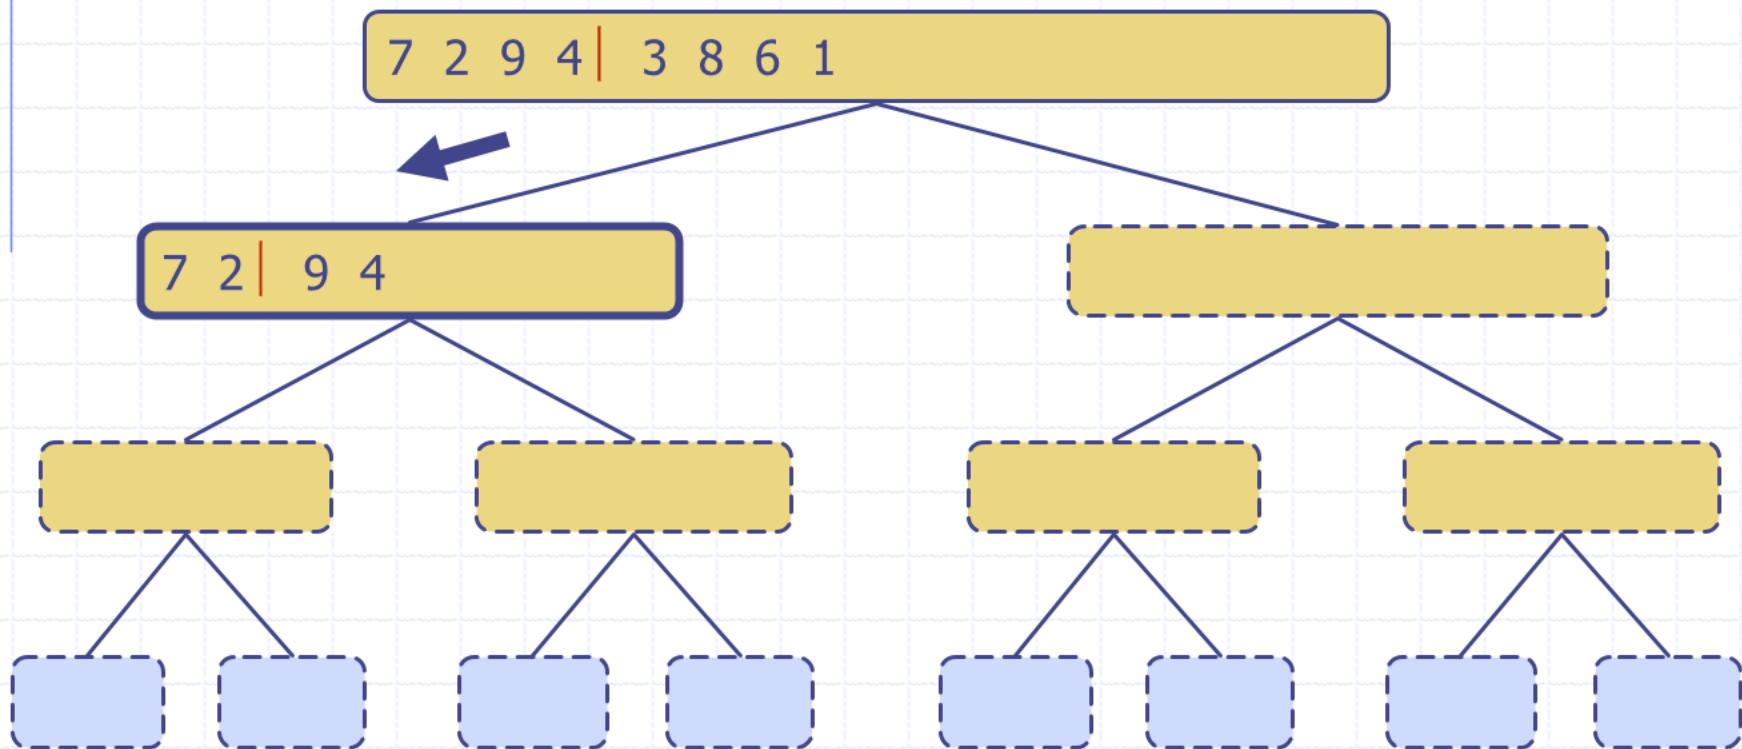
Execution Example

◆ Partition



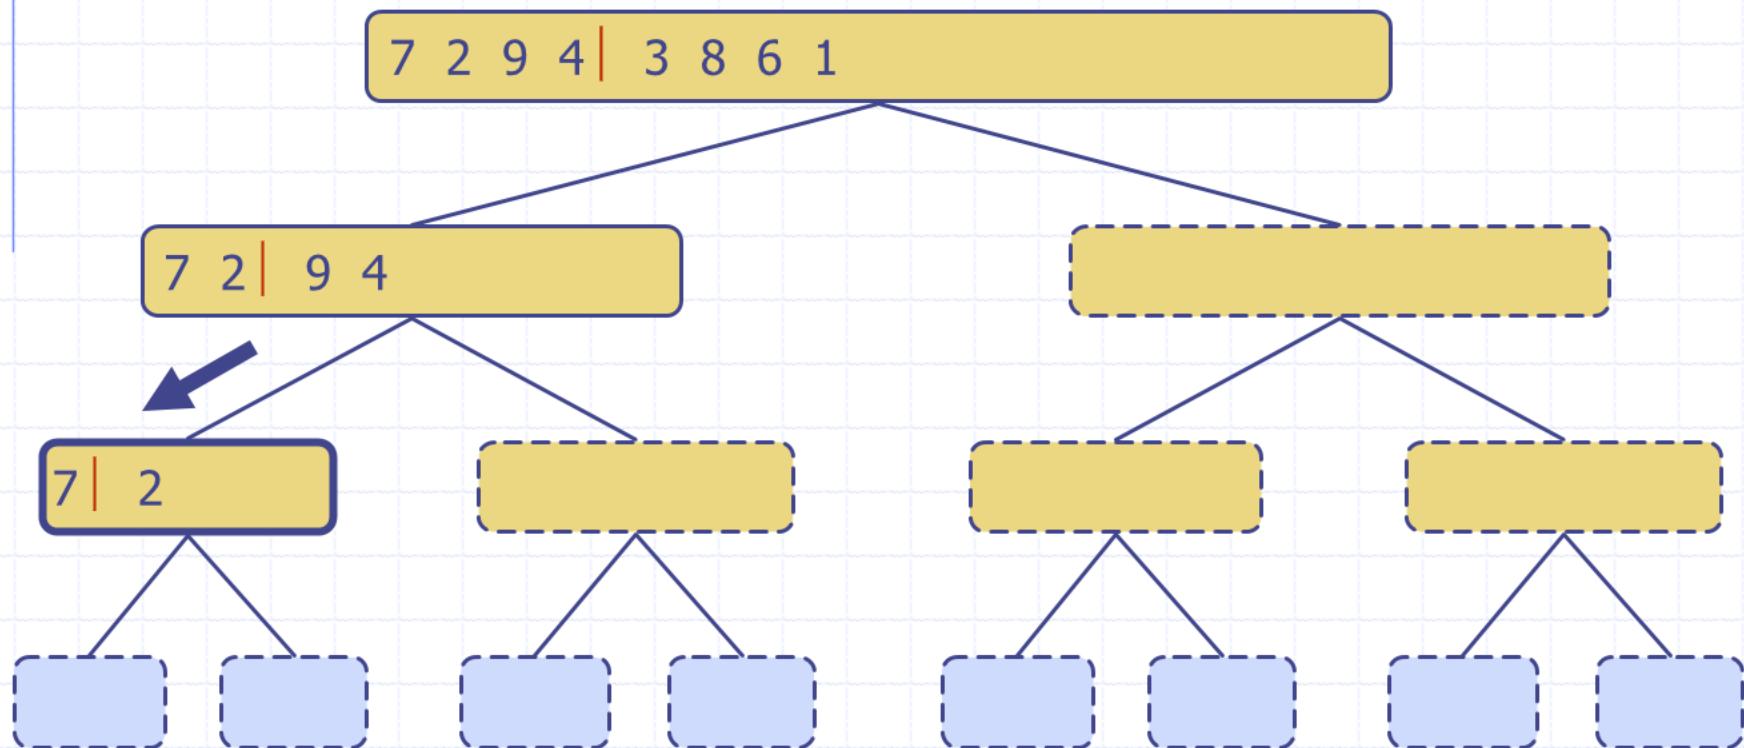
Execution Example (cont.)

- ◆ Recursive call, partition



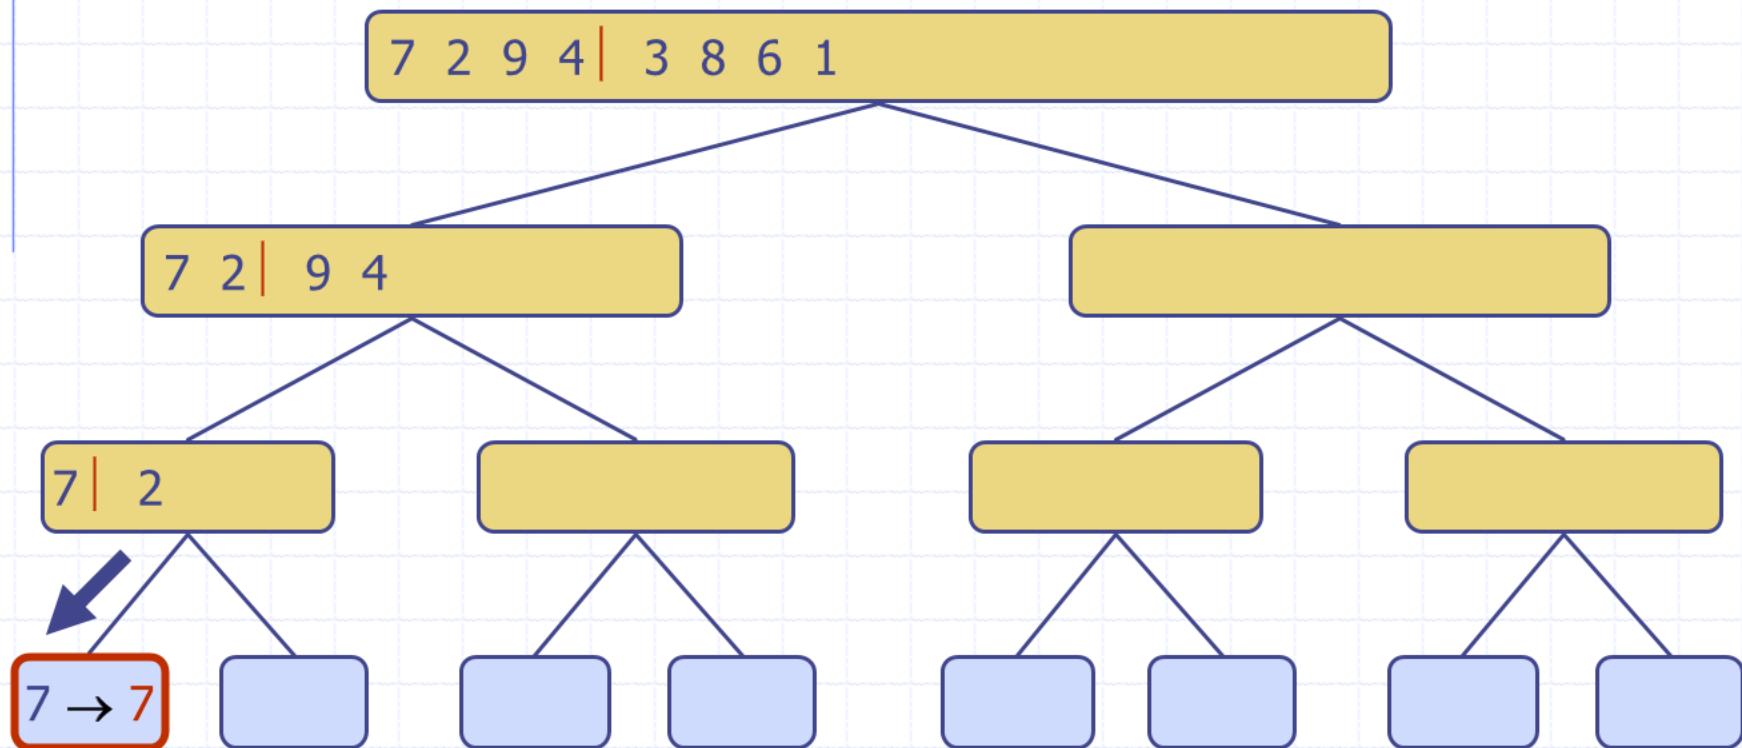
Execution Example (cont.)

- ◆ Recursive call, partition



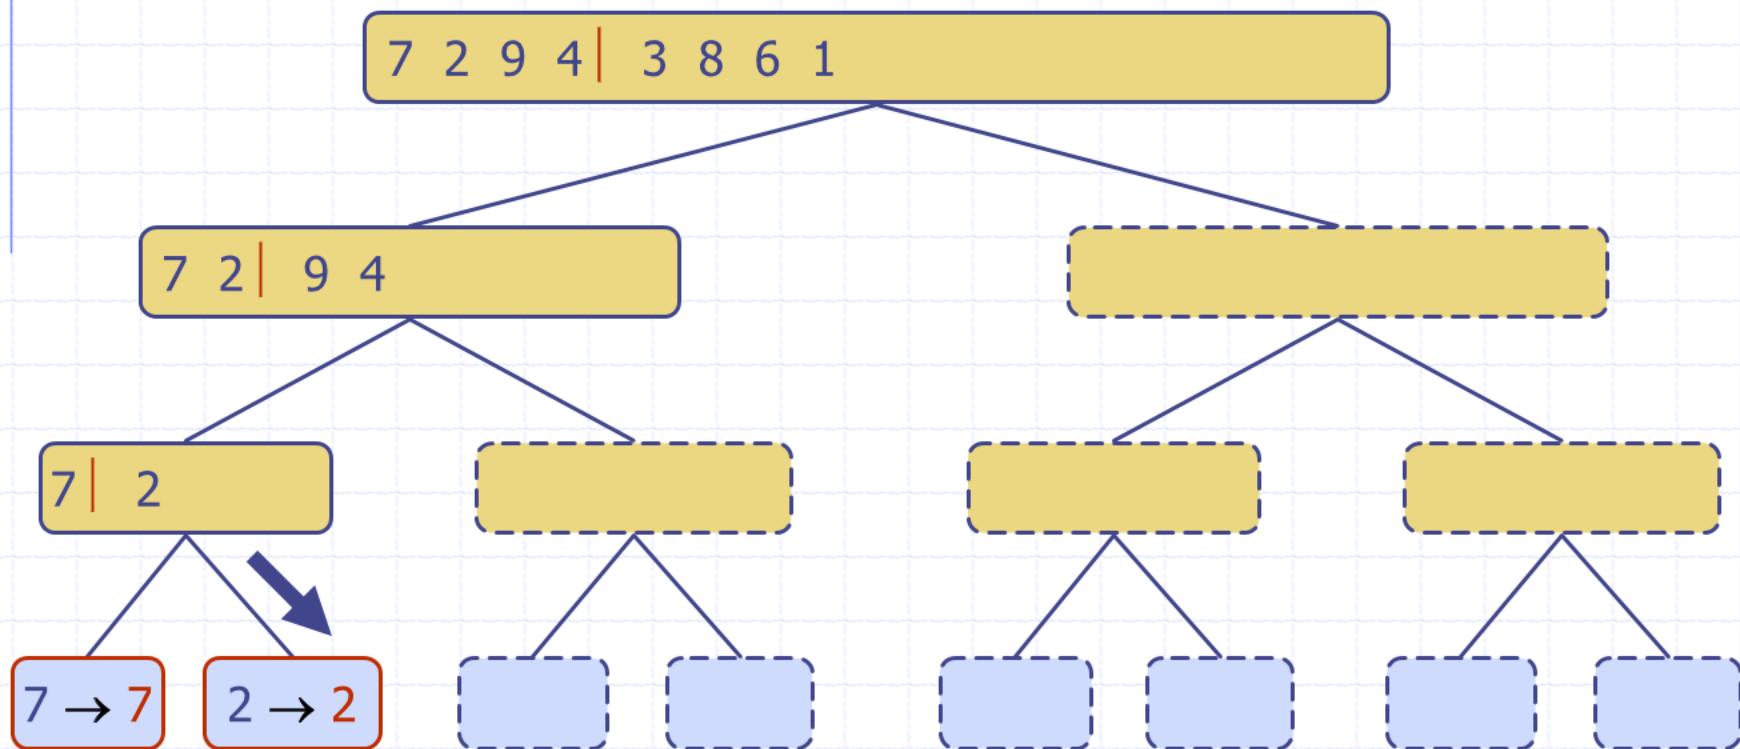
Execution Example (cont.)

- ◆ Recursive call, base case



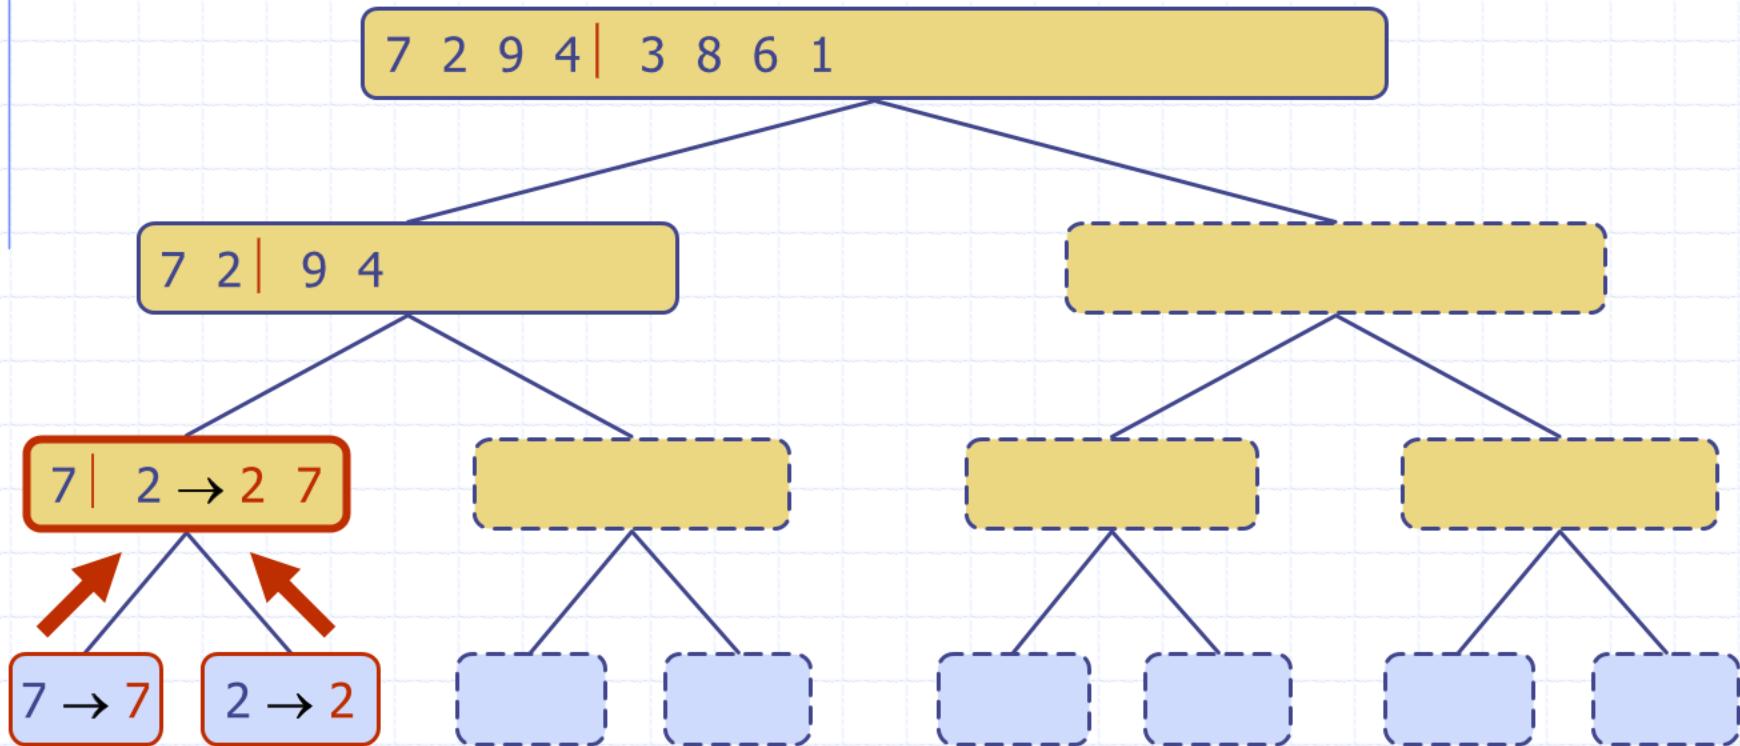
Execution Example (cont.)

- ◆ Recursive call, base case



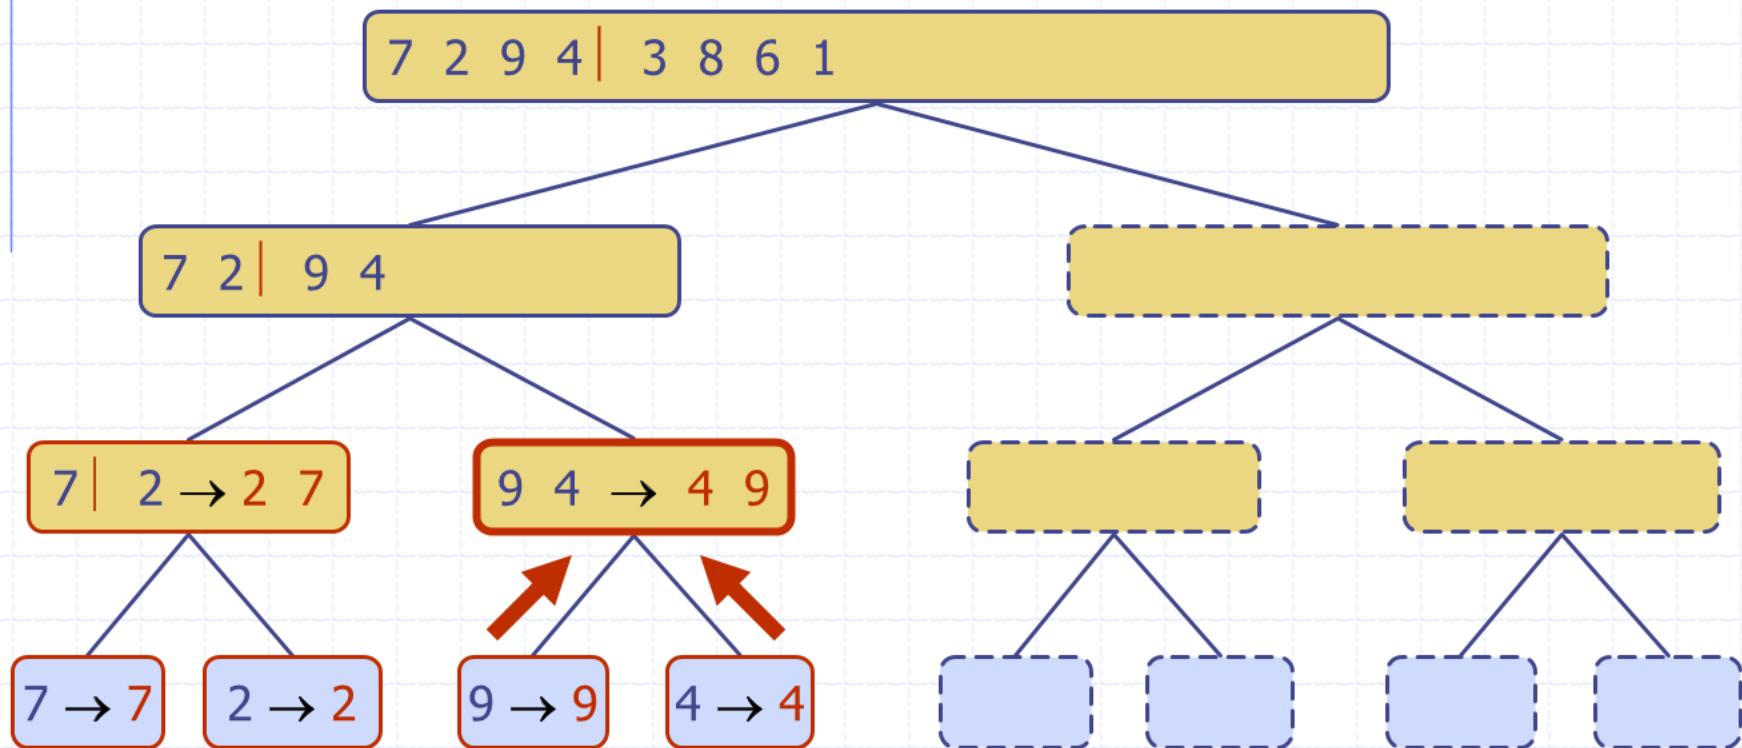
Execution Example (cont.)

◆ Merge



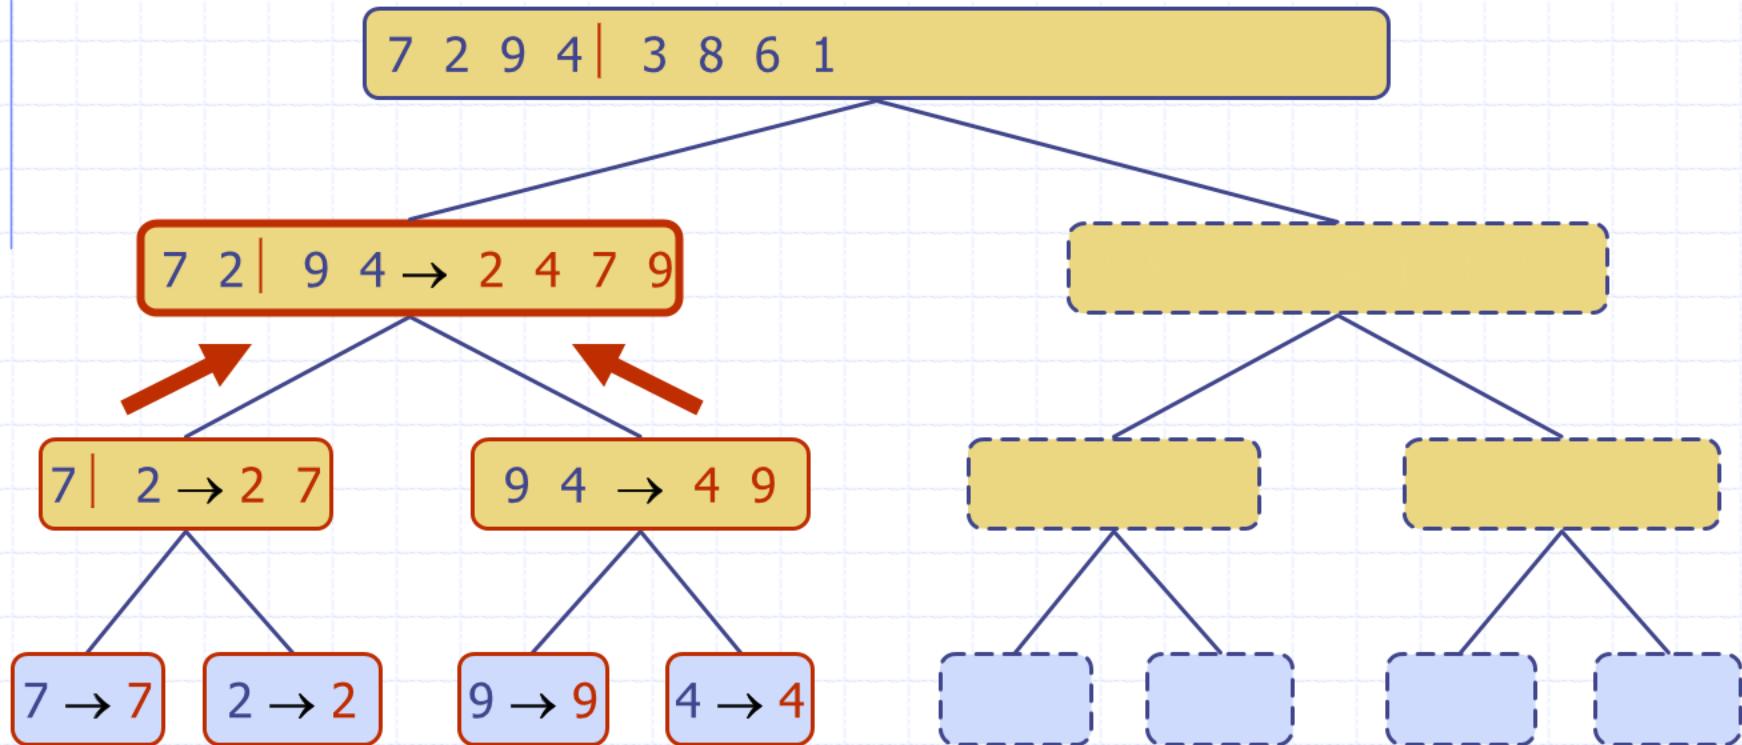
Execution Example (cont.)

- ◆ Recursive call, ..., base case, merge



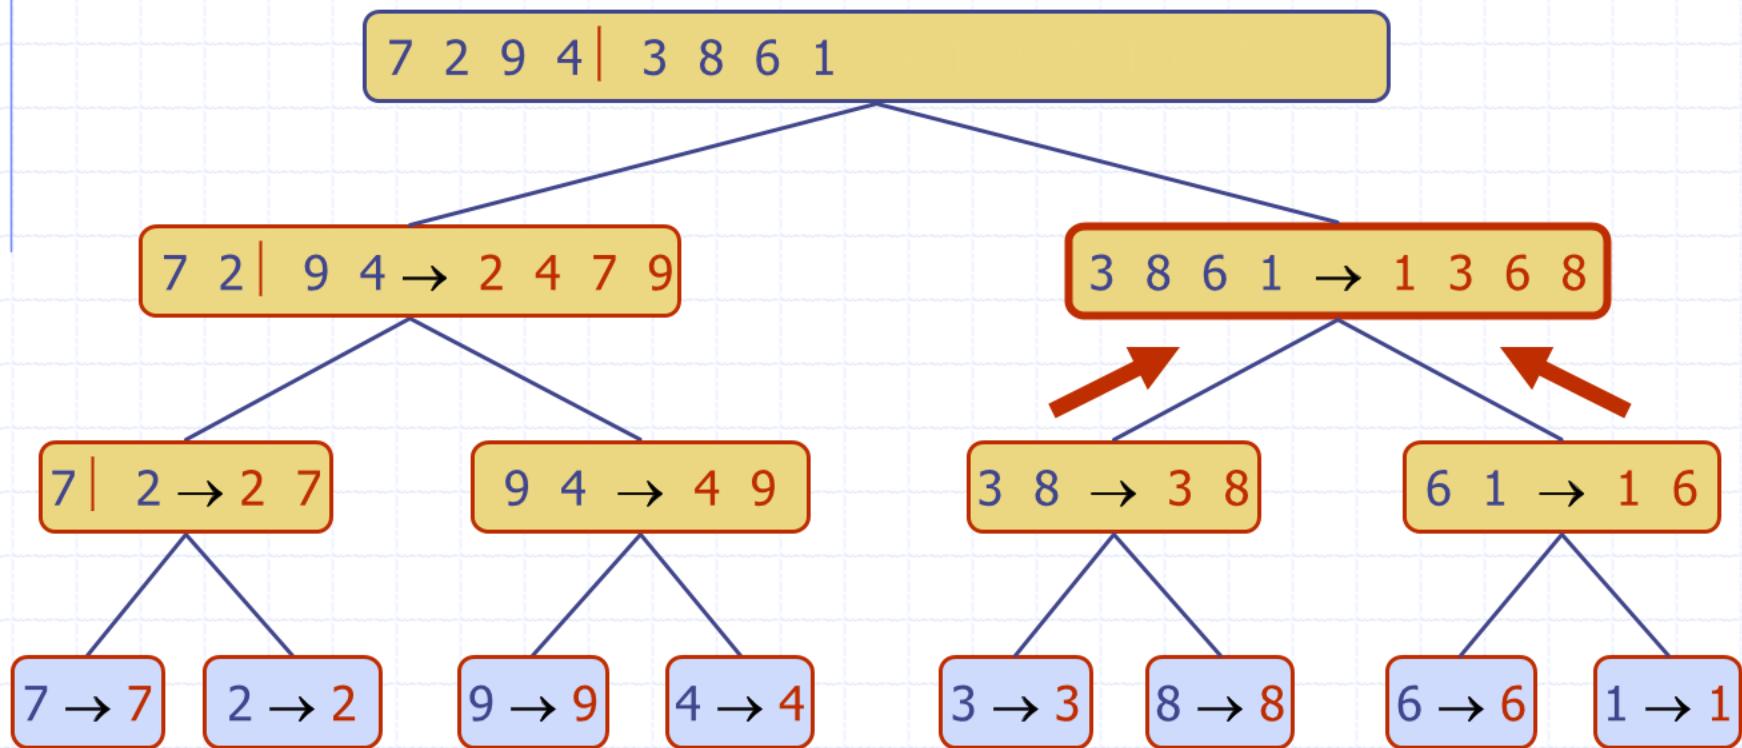
Execution Example (cont.)

◆ Merge



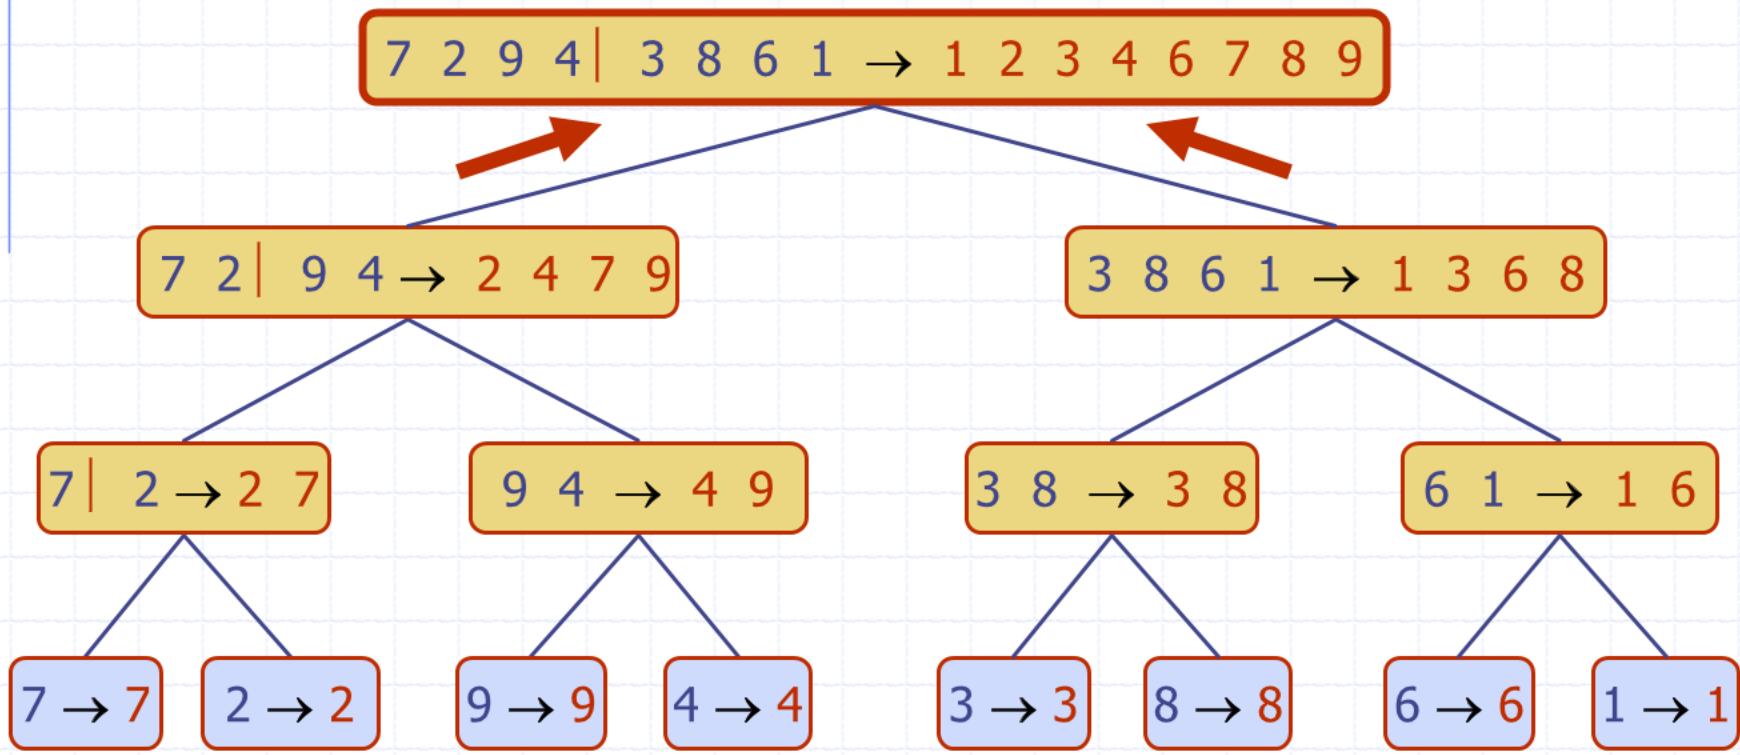
Execution Example (cont.)

- ◆ Recursive call, ..., merge, merge



Execution Example (cont.)

◆ Merge



Analysis of Merge-Sort

- ◆ The height h of the merge-sort tree is $O(\log n)$
 - at each recursive call we divide in half the sequence,
- ◆ The overall amount or work done at the nodes of depth i is $O(n)$
 - we partition and merge 2^i sequences of size $n/2^i$
 - we make 2^{i+1} recursive calls
- ◆ Thus, the total running time of merge-sort is $O(n \log n)$

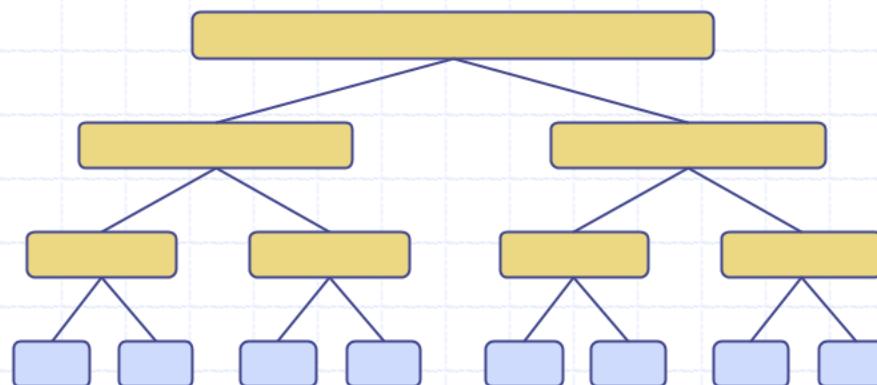
depth #seqs size

0 1 n

1 2 $n/2$

i 2^i $n/2^i$

... ...



Merge Sort

16

© 2004 Goodrich, Tamassia

$$\text{Time} = 1 \times n + 2 \times \frac{n}{2} + \dots + 2^i \times \frac{n}{2^i} = \sum_{i=0}^{\lfloor \log_2 n \rfloor} 2^i \cdot \frac{n}{2^i} = n \log_2 n$$

Another way of Solving:

$$\begin{aligned} T(n) &= 2 * T(n/2) + n = 4 * T(n/4) \\ &= 8 T(n/8) + 3n \dots = 2^n T(1) + h * n \\ \text{Sln as we saw is } T(n) &= n \log n + n \\ &\quad (\text{Verify}) \\ &= O(n \log n) \end{aligned}$$

Summary of Sorting Algorithms

Algorithm	Time	Notes
selection-sort	$O(n^2)$	◆ slow ◆ in-place ◆ for small data sets (< 1K)
insertion-sort	$O(n^2)$	◆ slow ◆ in-place ◆ for small data sets (< 1K)
heap-sort	$O(n \log n)$	◆ fast ◆ in-place ◆ for large data sets (1K — 1M)
merge-sort	$O(n \log n)$	◆ fast ◆ sequential data access ◆ for huge data sets (> 1M)

{ By you need additional array space } design

Nonrecursive Merge-Sort

merge runs of length 2, then 4, then 8, and so on

merge two runs in the in array to the out array

```
public static void mergeSort(Object[] orig, Comparator c) { // nonrecursive
    Object[] in = new Object[orig.length]; // make a new temporary array
    System.arraycopy(orig,0,in,0,in.length); // copy the input
    Object[] out = new Object[in.length]; // output array
    Object[] temp; // temp array reference used for swapping
    int n = in.length;
    for (int i=1; i < n; i*=2) { // each iteration sorts all length-2*i runs
        for (int j=0; j < n; j+=2*i) // each iteration merges two length-i pairs
            merge(in,out,c,j,i); // merge from in to out two length-i runs at j
        temp = in; in = out; out = temp; // swap arrays for next iteration
    }
    // the "in" array contains the sorted array, so re-copy it
    System.arraycopy(in,0,orig,0,in.length);
}
protected static void merge(Object[] in, Object[] out, Comparator c, int start,
    int inc) { // merge in[start..start+inc-1] and in[start+inc..start+2*inc-1]
    int x = start; // index into run #1
    int end1 = Math.min(start+inc, in.length); // boundary for run #1
    int end2 = Math.min(start+2*inc, in.length); // boundary for run #2
    int y = start+inc; // index into run #2 (could be beyond array boundary)
    int z = start; // index into the out array
    while ((x < end1) && (y < end2))
        if (c.compare(in[x],in[y]) <= 0) out[z++] = in[x++];
        else out[z++] = in[y++];
    if (x < end1) // first run didn't finish
        System.arraycopy(in, x, out, z, end1 - x);
    else if (y < end2) // second run didn't finish
        System.arraycopy(in, y, out, z, end2 - y);
}
```

Merge Sort

Some next questions:

① Does division have to be "balanced" at every step?

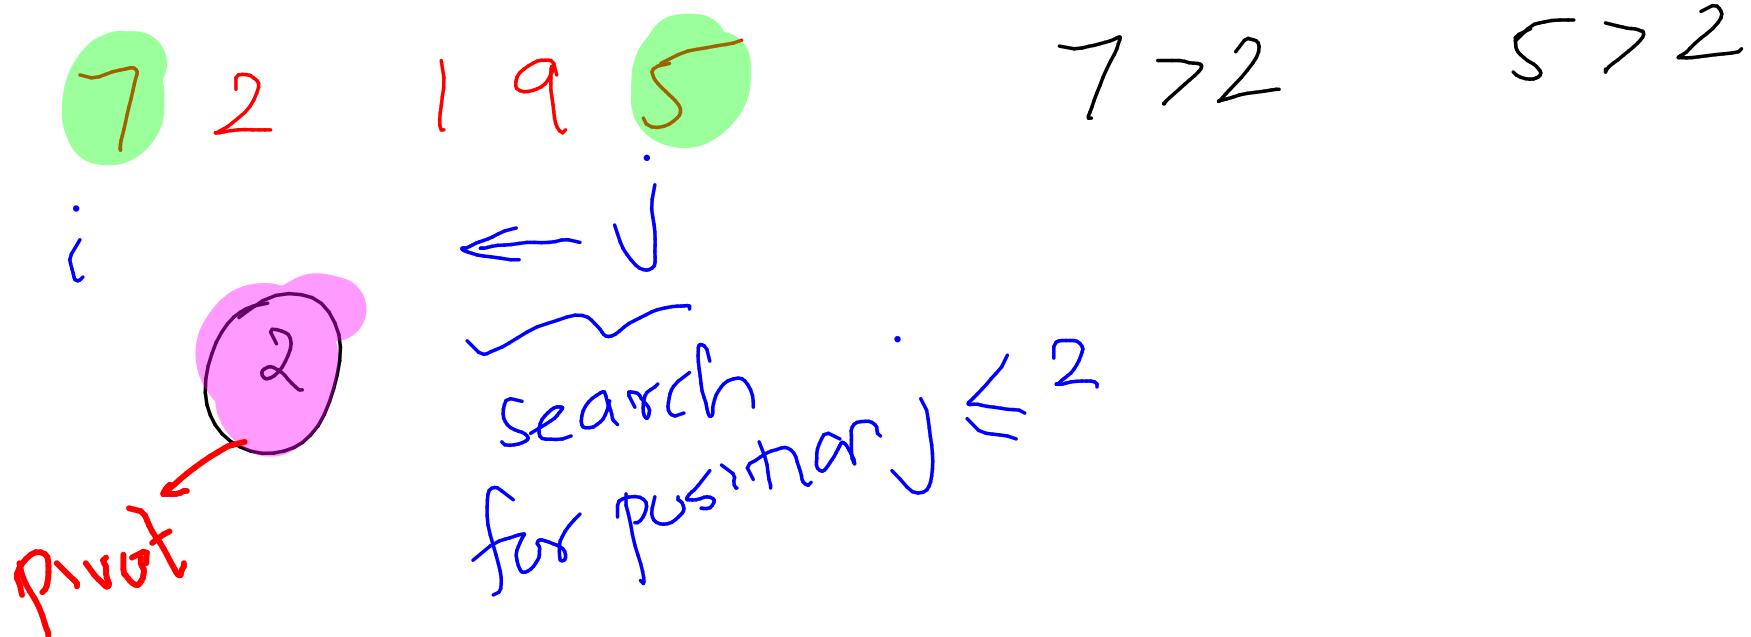
ie $n \rightarrow \frac{n}{2}, \frac{n}{2}$ is necessary?

② Should the division be based on length only? Can it be based on "pivot"

Q: Can I use pivot to do in-place divide?

↓ pivot element

new pivots



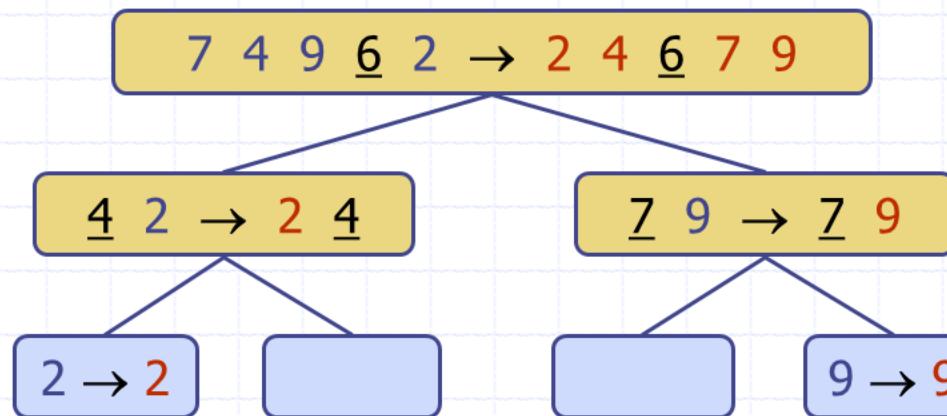
Idea: Move i right if element at j is $<$ pivot & at i is $<$ pivot until element at i becomes $>$ pivot.

Swap elements at i & j

Similarly: Move j left if element at i is $>$ pivot & at j is $>$ pivot until --

Quick-Sort

Divide
on based
on pivot



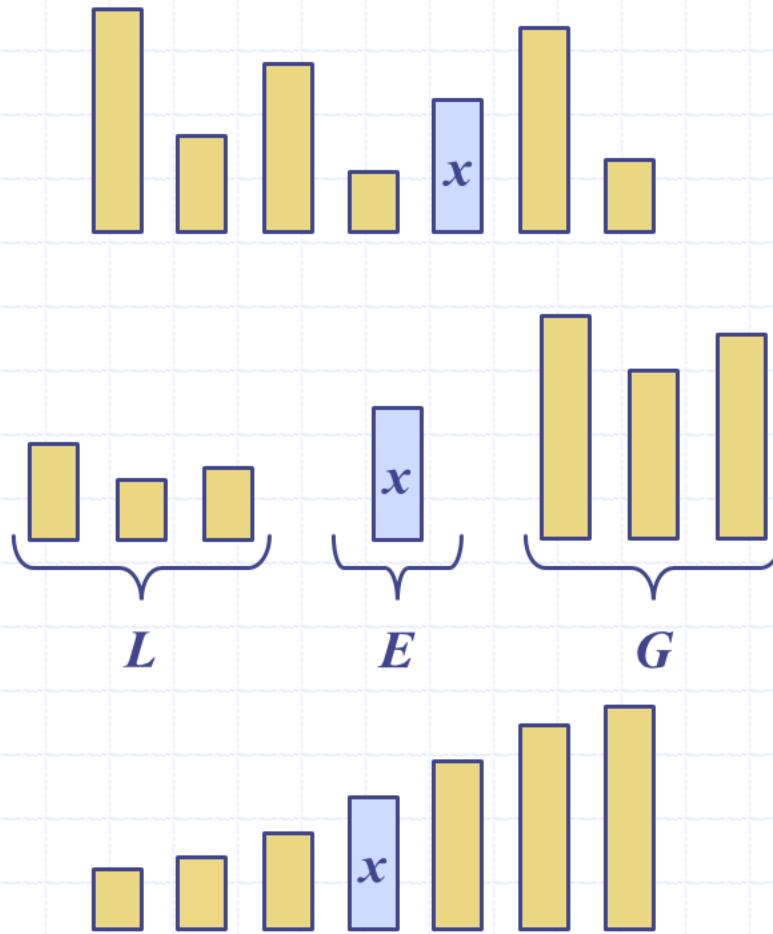
Merge
is simply
concatenat
-ing

You
even need
step when you divide
not
1 a merge
in-place

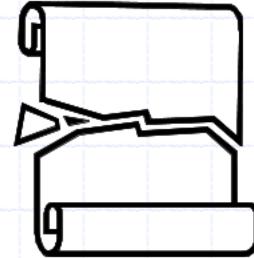
Quick-Sort (§ 10.2)

◆ Quick-sort is a randomized sorting algorithm based on the divide-and-conquer paradigm:

- Divide: pick a random element x (called pivot) and partition S into
 - L elements less than x
 - E elements equal x
 - G elements greater than x
- Recur: sort L and G
- Conquer: join L , E and G



Partition



in place
version

- ◆ We partition an input sequence as follows:
 - We remove, in turn, each element y from S and
 - We insert y into L , E or G , depending on the result of the comparison with the pivot x
- ◆ Each insertion and removal is at the beginning or at the end of a sequence, and hence takes $O(1)$ time
- ◆ Thus, the partition step of quick-sort takes $O(n)$ time

Algorithm *partition(S, p)*

Input sequence S , position p of pivot

Output subsequences L , E , G of the elements of S less than, equal to, or greater than the pivot, resp.

$L, E, G \leftarrow$ empty sequences

$x \leftarrow S.\text{remove}(p)$

while $\neg S.\text{isEmpty}()$

$y \leftarrow S.\text{remove}(S.\text{first}())$

if $y < x$

$L.\text{insertLast}(y)$

else if $y = x$

$E.\text{insertLast}(y)$

else { $y > x$ }

$G.\text{insertLast}(y)$

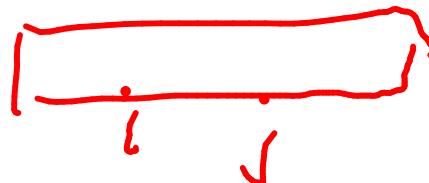
return L, E, G

This is not in place. Better to use the in place version we discussed

Array: $A[1..n]$

e = pivot element

i=1, j=n



if ($A[i] > e$) && ($A[j] > e$) j--

if ($A[i] \leq e$) && ($A[j] < e$) i++

if ($A[i] > e$) && ($A[j] \leq e$) { swap $A[i]$ & $A[j]$ }

i++
j--

}

In-Place Quick-Sort



- ◆ Quick-sort can be implemented to run in-place
- ◆ In the partition step, we use replace operations to rearrange the elements of the input sequence such that
 - the elements less than the pivot have rank less than h
 - the elements equal to the pivot have rank between h and k
 - the elements greater than the pivot have rank greater than k
- ◆ The recursive calls consider
 - elements with rank less than h
 - elements with rank greater than k

Algorithm *inPlaceQuickSort(S, l, r)*

Input sequence S , ranks l and r

Output sequence S with the elements of rank between l and r rearranged in increasing order

if $l \geq r$

return

$i \leftarrow$ a random integer between l and r

$x \leftarrow S.elemAtRank(i)$

$(h, k) \leftarrow inPlacePartition(x)$

$inPlaceQuickSort(S, l, h - 1)$

$inPlaceQuickSort(S, k + 1, r)$

In-Place Partitioning



- ◆ Perform the partition using two indices to split S into L and E U G (a similar method can split E U G into E and G).

j

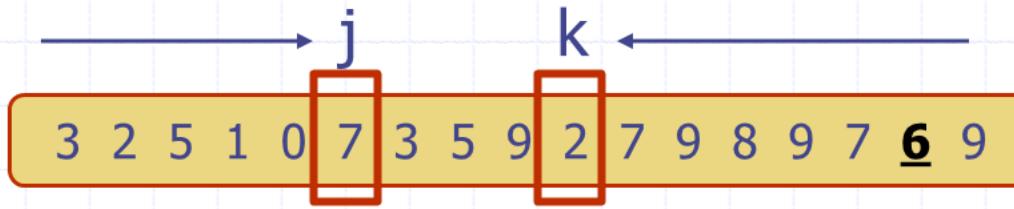
k

3 2 5 1 0 7 3 5 9 2 7 9 8 9 7 **6** 9

(pivot = 6)

- ◆ Repeat until j and k cross:

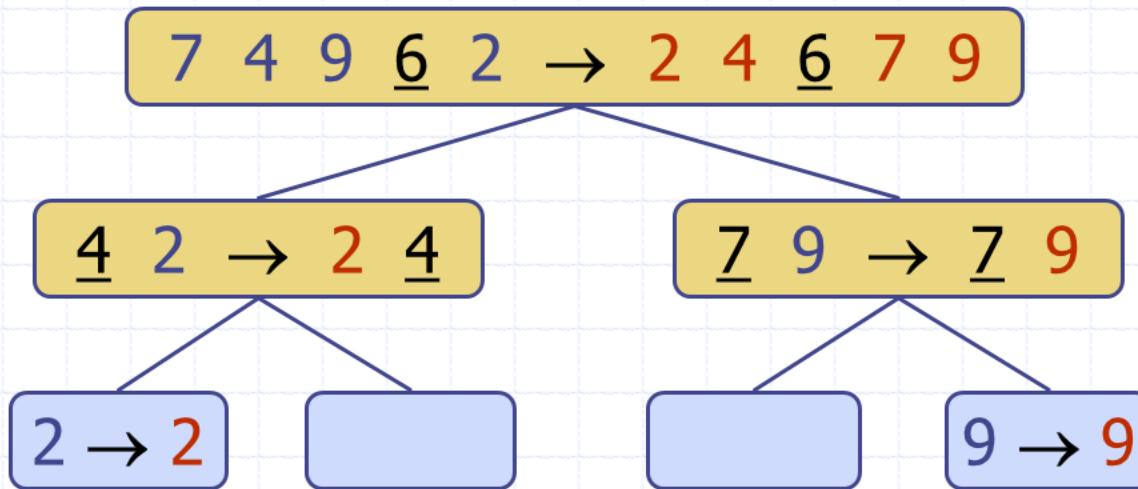
- Scan j to the right until finding an element $\geq x$.
- Scan k to the left until finding an element $< x$.
- Swap elements at indices j and k



Quick-Sort

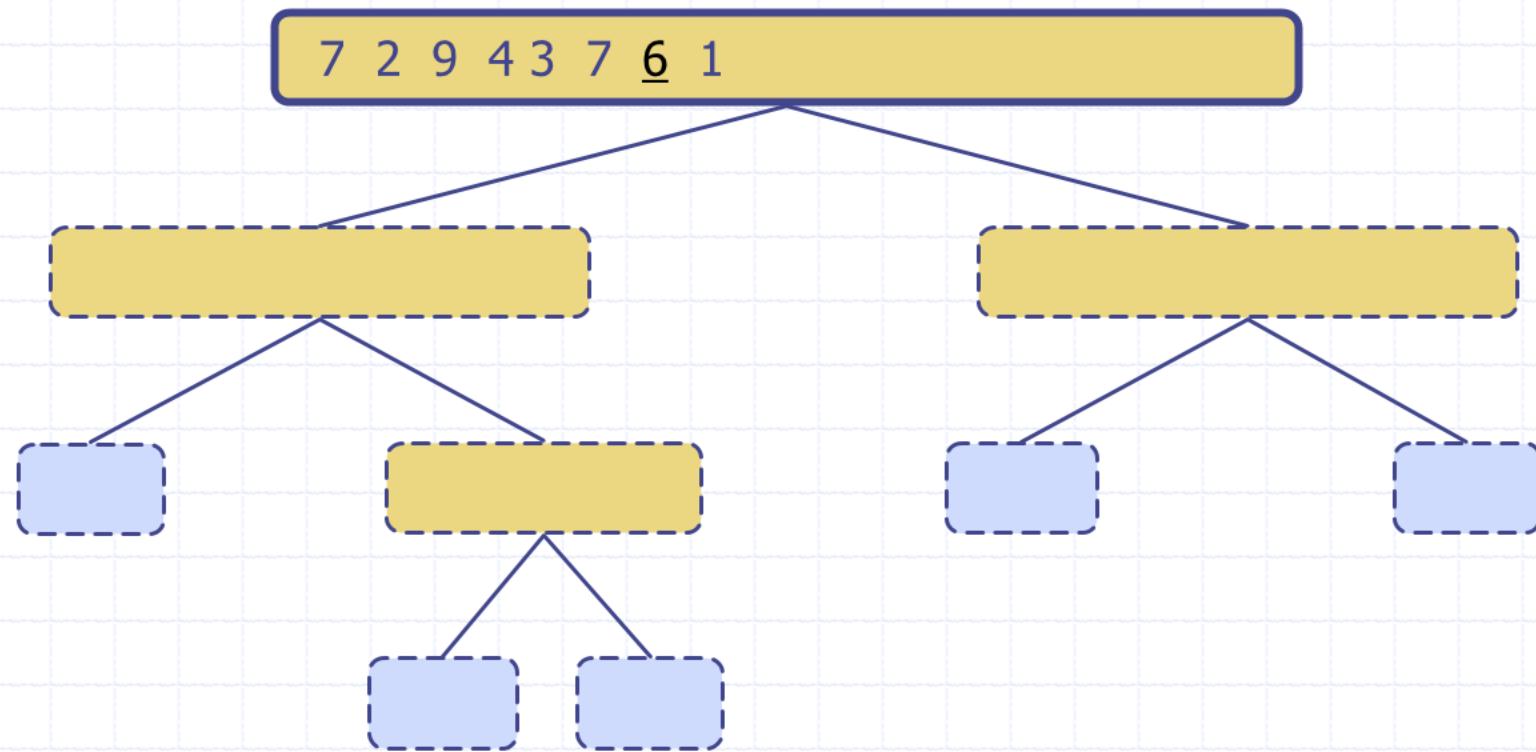
Quick-Sort Tree

- ◆ An execution of quick-sort is depicted by a binary tree
 - Each node represents a recursive call of quick-sort and stores
 - ◆ Unsorted sequence before the execution and its pivot
 - ◆ Sorted sequence at the end of the execution
 - The root is the initial call
 - The leaves are calls on subsequences of size 0 or 1



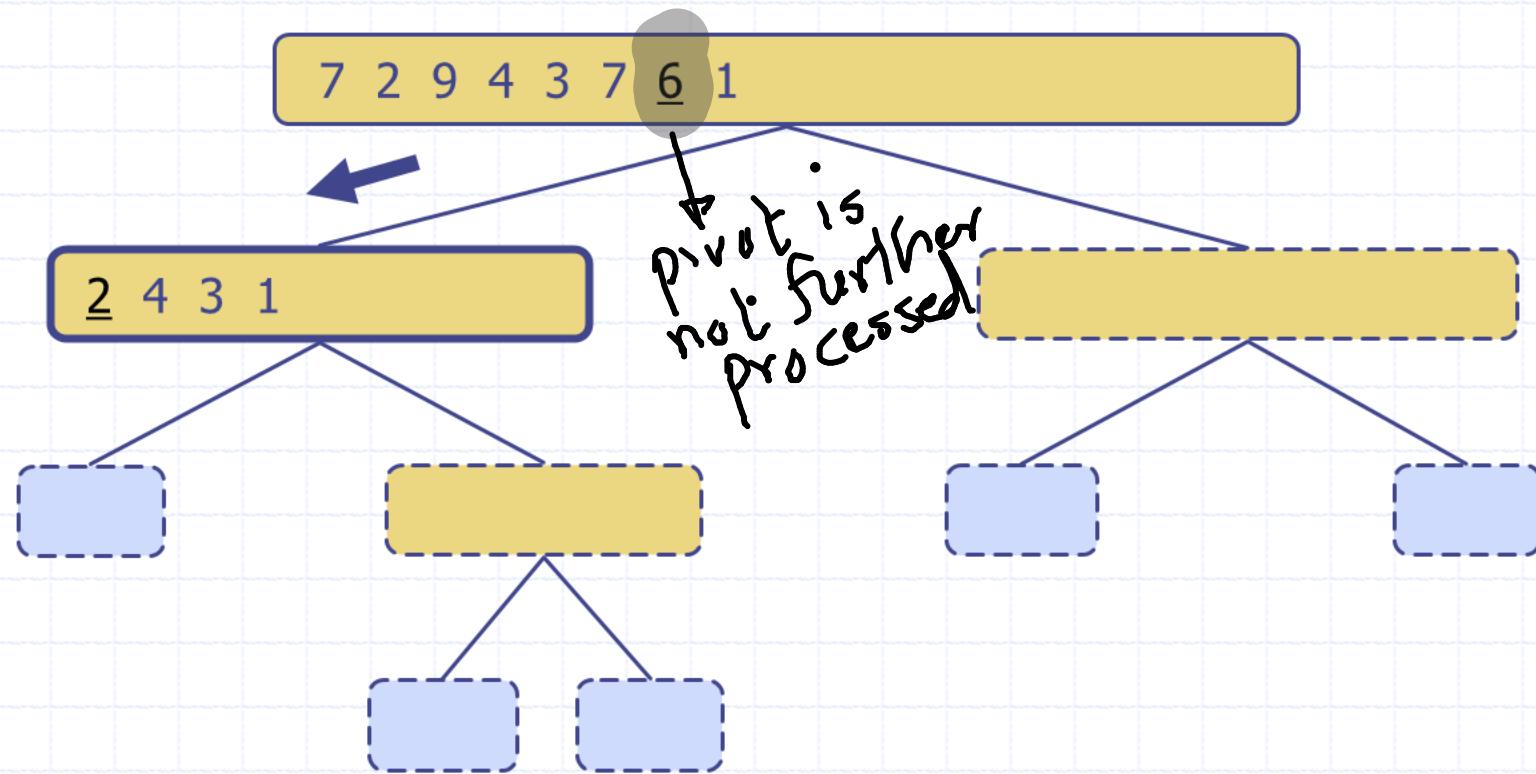
Execution Example

◆ Pivot selection



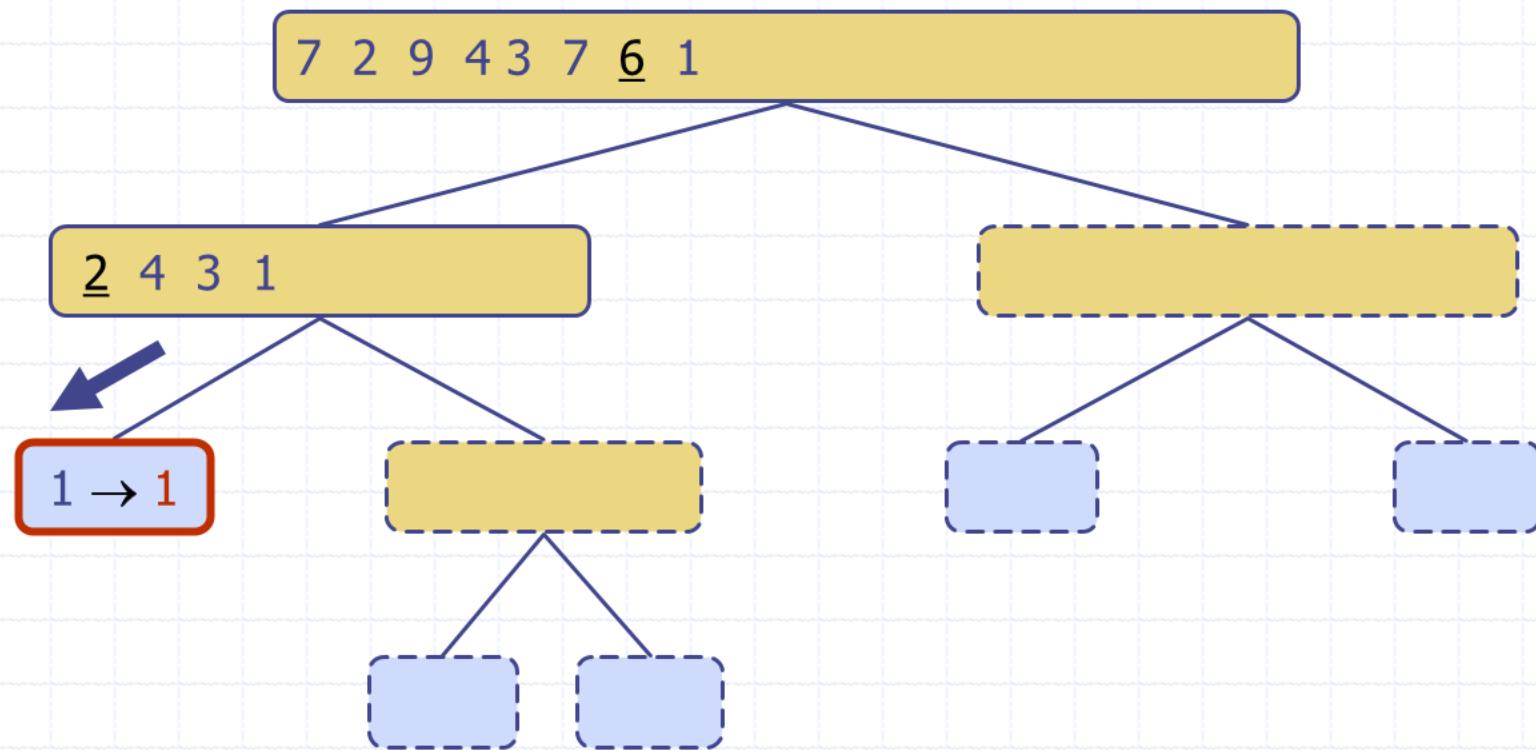
Execution Example (cont.)

- ◆ Partition, recursive call, pivot selection



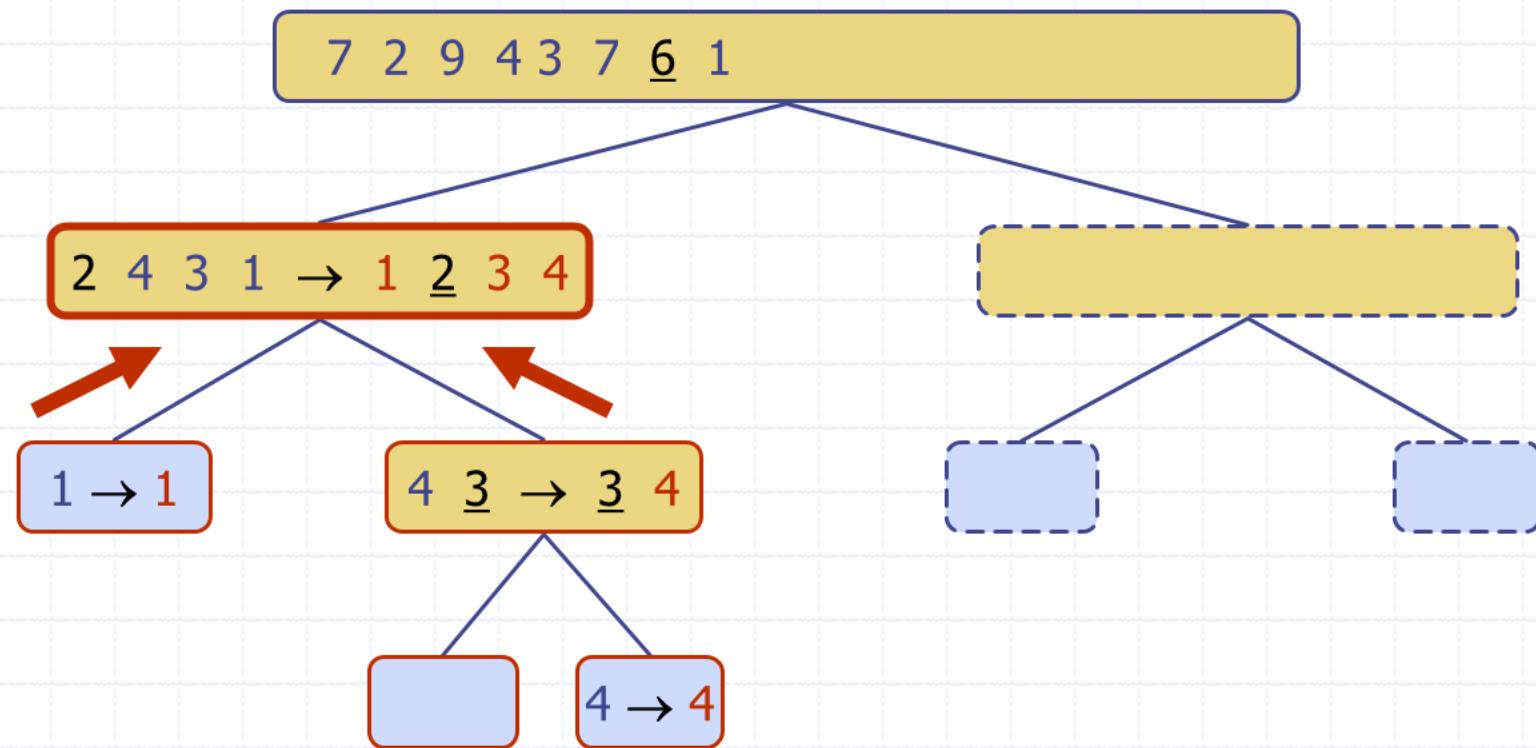
Execution Example (cont.)

- ◆ Partition, recursive call, base case



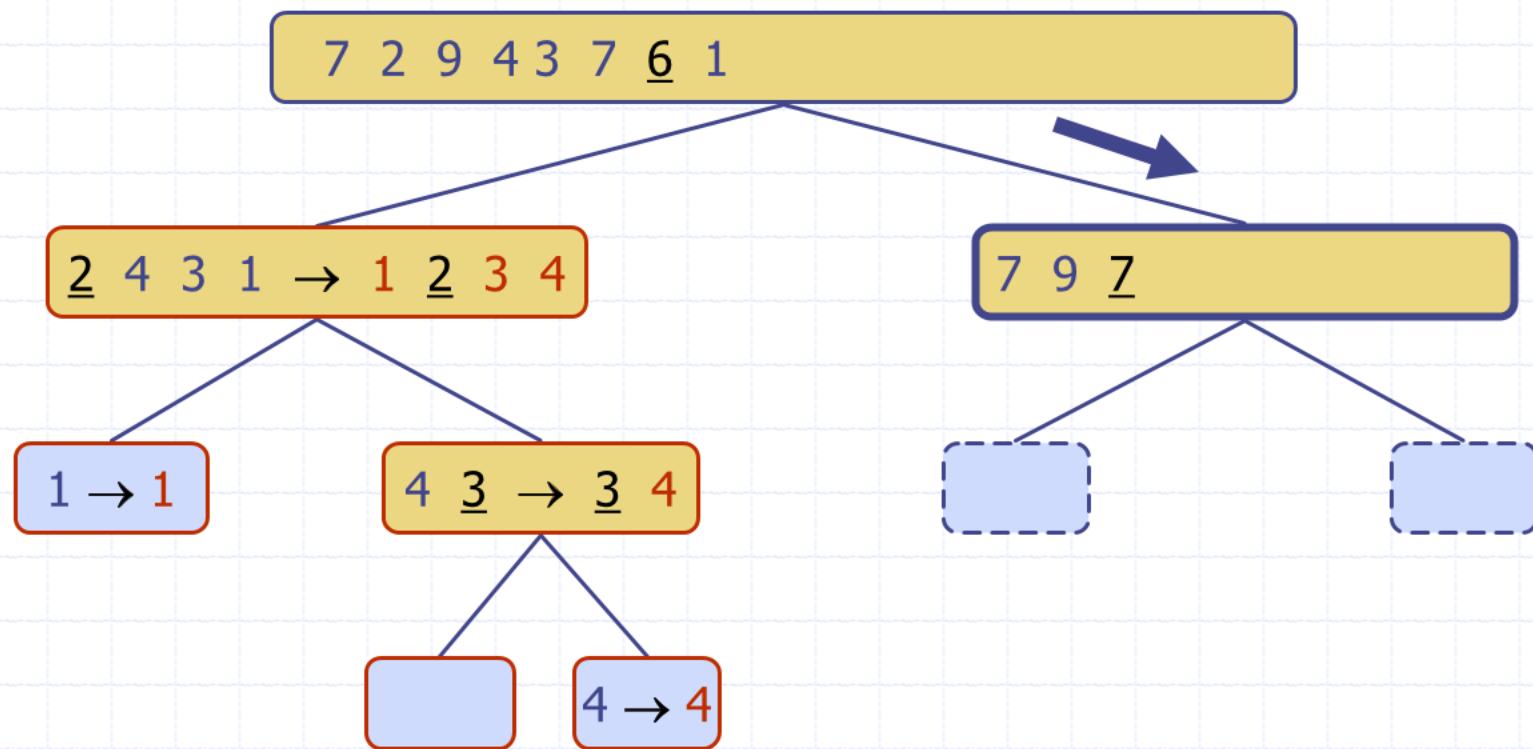
Execution Example (cont.)

- ◆ Recursive call, ..., base case, join



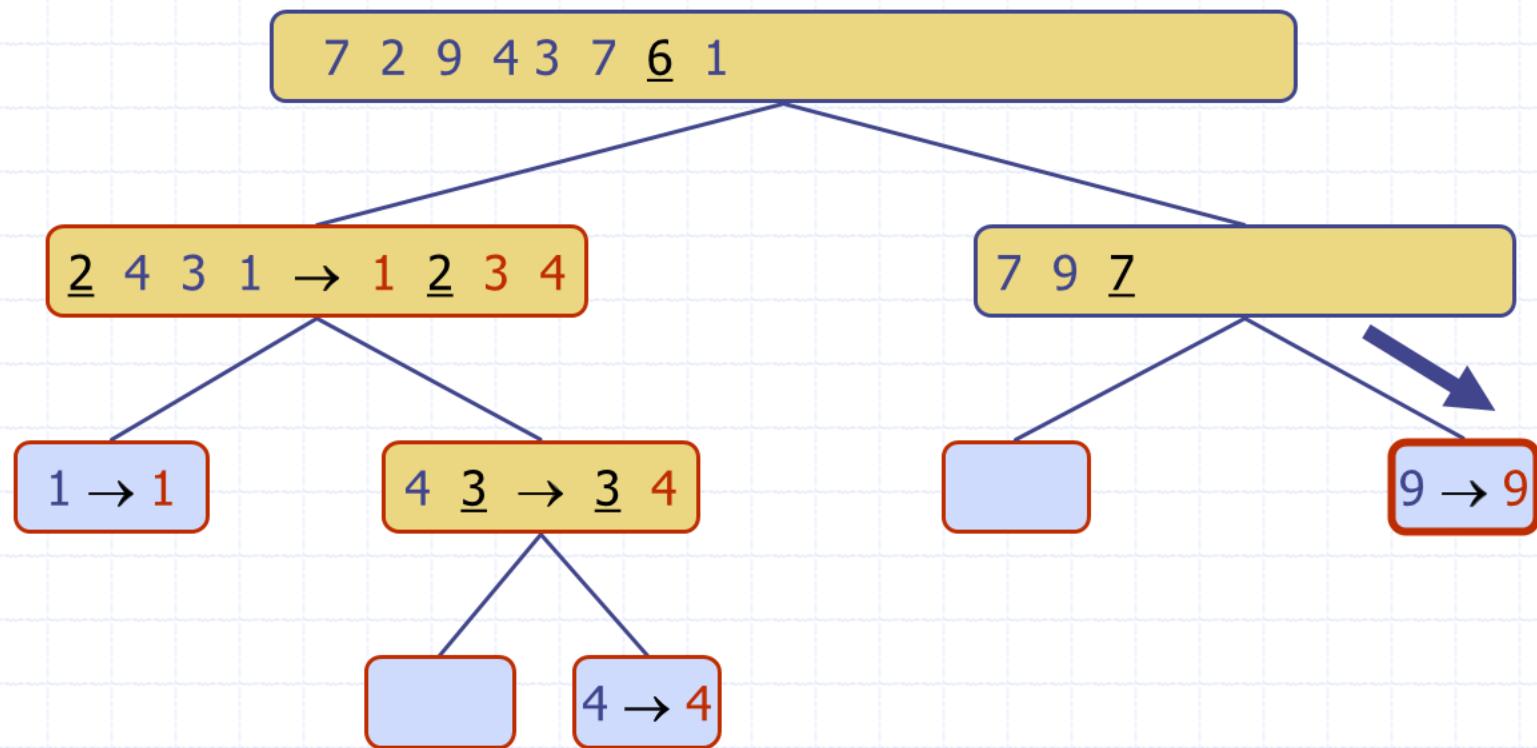
Execution Example (cont.)

- ◆ Recursive call, pivot selection



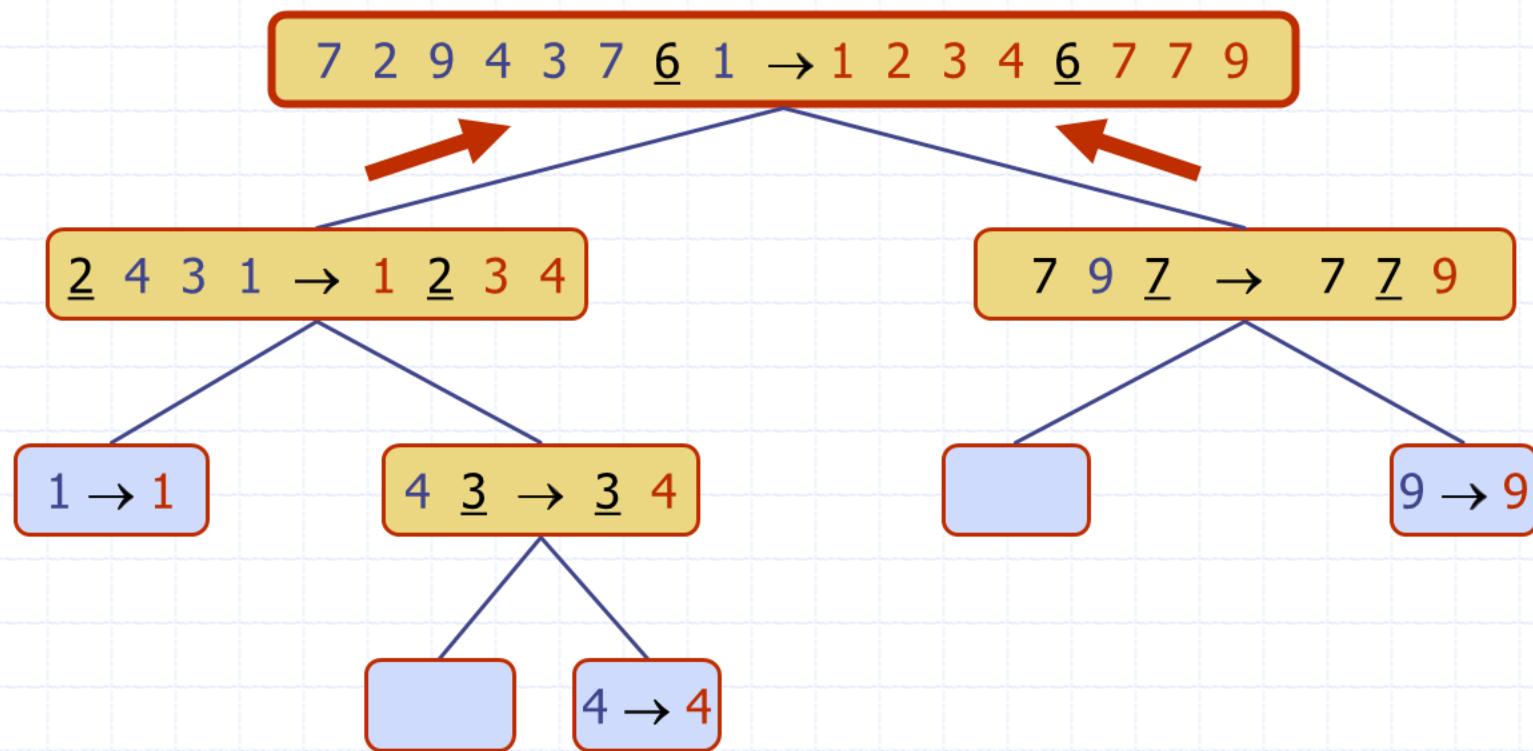
Execution Example (cont.)

◆ Partition, ..., recursive call, base case



Execution Example (cont.)

Join, join



Worst-case Running Time

- ◆ The worst case for quick-sort occurs when the pivot is the unique minimum or maximum element
- ◆ One of L and G has size $n - 1$ and the other has size 0
- ◆ The running time is proportional to the sum
$$n + (n - 1) + \dots + 2 + 1 = O(n^2)$$
- ◆ Thus, the worst-case running time of quick-sort is $O(n^2)$

